

# Intelligence System: Reinforcement Learning in Machine Learning



សាកលវិទ្យាល័យន័រតុន  
NORTON UNIVERSITY

Intelligence System  
Development

2024 – 2025  
Y4E1 – DCS – NU

**By: SEK SOCHEAT**

Advisor to DCS and Lecturer

**Mobile:** 017 879 967

**Email:** [socheat.sek@gmail.com](mailto:socheat.sek@gmail.com)

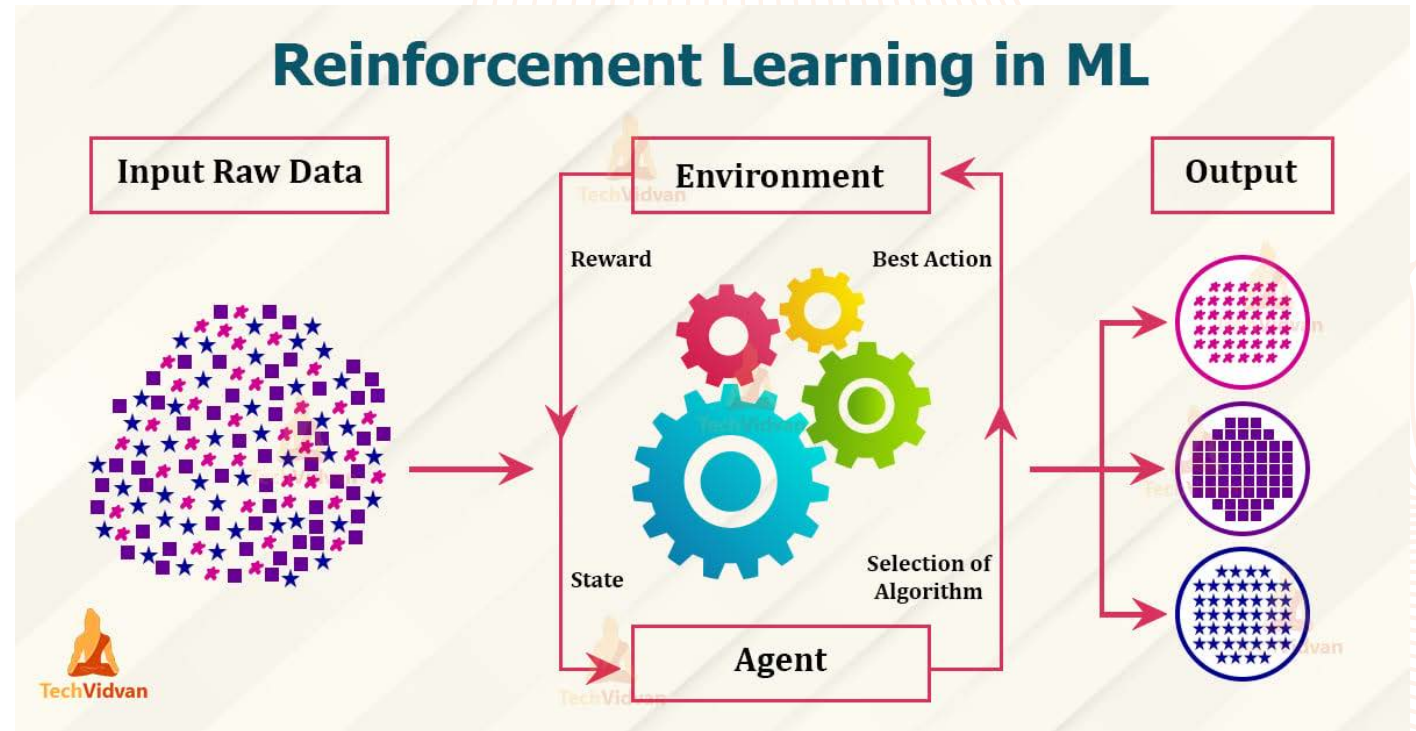
# Table of Contents

- **Introduction to Reinforcement Learning**
- **Key Characteristics of Reinforcement Learning**
- **Types of Reinforcement Learning**
- **Popular Algorithms in Reinforcement Learning**
- **Steps in Reinforcement Learning**
- **Applications of Reinforcement Learning**
- **Challenges in Reinforcement Learning**
- **Example**
- **Homework**

# Introduction to Reinforcement Learning

Reinforcement learning (RL) is a field within machine learning focused on how agents can learn optimal behaviors through interactions with their environment.

The goal is to maximize cumulative rewards over time. It emphasizes decision-making, where an agent learns by receiving feedback in the form of rewards or penalties in response to its actions, rather than being explicitly told what to do.



# Key Characteristics of Reinforcement Learning

- 1. Agent-Environment Interaction:** The agent interacts with the environment, taking actions and receiving feedback in the form of rewards.
- 2. Trial and Error Learning:** RL relies on learning through exploration (trying new actions) and exploitation (using known actions).
- 3. Delayed Rewards:** Actions may lead to long-term benefits rather than immediate rewards, requiring planning.
- 4. Policy:** A strategy or mapping from states to actions that the agent learns to maximize rewards.
- 5. Value Functions:** These estimate the future rewards from a given state or action, guiding the agent's decisions.



# Key Characteristics of Reinforcement Learning

- 6. Exploration vs. Exploitation:** A balance is needed between exploring unknown actions and exploiting known ones for rewards.
- 7. Markov Decision Process (MDP):** RL often assumes the problem can be modeled as an MDP, with defined states, actions, rewards, and transitions.
- 8. No Supervised Labels:** Unlike supervised learning, RL does not rely on labeled input-output pairs but learns through reward signals.
- 9. Feedback-driven:** The agent learns dynamically from rewards and penalties, improving iteratively.

# Types of Reinforcement Learning

Reinforcement Learning (RL) can be categorized into two main types based on how the agent learns to interact with its environment:

*1. Model-Based Reinforcement Learning*

*2. Model-Free Reinforcement Learning*

# Types of Reinforcement Learning

## *1. Model-Based Reinforcement Learning*

- **Definition:** The agent builds a model of the environment, including the transition probabilities and reward structure.
- **Purpose:** Enables planning by simulating potential future states and actions.
- **Advantages:** Efficient in solving problems where a model can be accurately learned; allows leveraging planning algorithms like dynamic programming.
- **Challenges:** Difficult when the environment is complex or uncertain.

# Types of Reinforcement Learning

## *2. Model-Free Reinforcement Learning*

- **Definition:** The agent learns directly from interactions without building a model of the environment.
- **Purpose:** Focuses on learning the optimal policy or value functions.
- **Common Algorithms:**
  - *Q-Learning:* Learns the value of state-action pairs to guide decision-making.
  - *SARSA:* Learns the value of taking specific actions under the current policy.
- **Advantages:** Simpler to implement and often effective for complex environments.
- **Challenges:** May require more interactions with the environment, leading to inefficiency in some cases.



# Popular Algorithms in Reinforcement Learning

## *1. Model-Free Algorithms*

### *Q-Learning*


- Off-policy algorithm.
- Learns the optimal action-value function without requiring a model of the environment.
- Updates are based on maximum estimated future rewards.

### *SARSA (State-Action-Reward-State-Action)*

- On-policy algorithm.
- Updates based on the actual actions taken by the current policy, making it more sensitive to the policy's behavior.

### *Deep Q-Networks (DQN)*

- Combines Q-learning with deep neural networks to handle high-dimensional state spaces.
- Introduces experience replay and target networks for stability.



# Popular Algorithms in Reinforcement Learning

## *2. Policy Gradient Methods*

### ***REINFORCE***

- Directly learns the policy by maximizing the expected cumulative reward using gradient ascent.

### ***Actor-Critic***


- Combines policy learning (actor) and value function learning (critic) for stability and efficiency.

### ***Proximal Policy Optimization (PPO)***

- Simplifies and stabilizes policy updates, ensuring bounded changes during training.

### ***Trust Region Policy Optimization (TRPO)***

- Improves stability by constraining the policy updates within a trust region.



# Popular Algorithms in Reinforcement Learning


## *3. Model-Based Algorithms*

### *Dyna-Q*

- Combines learning and planning, using a model of the environment for faster convergence.

### *Monte Carlo Tree Search (MCTS)*

- Used in games like chess and Go, it builds a search tree of possible actions and rewards using a model.



# Popular Algorithms in Reinforcement Learning

## *4. Other Advanced Methods*

### *Temporal-Difference (TD) Learning*

- Combines ideas from Monte Carlo methods and dynamic programming to update values incrementally.


### *Double Q-Learning*

- Addresses overestimation bias in Q-learning by using two Q-functions for updates.

### *Asynchronous Advantage Actor-Critic (A3C)*

- Leverages multiple threads for faster and more stable learning.

These algorithms form the backbone of RL advancements, and their choice depends on the problem complexity and requirements.



# Steps in Reinforcement Learning

Reinforcement Learning (RL) typically follows these key steps:

## 1. Define the Environment:

- Specify the environment where the agent will interact.
- Define states, actions, rewards, and transitions (usually modeled as a Markov Decision Process).

## 2. Initialize the Agent:

- Set up the agent's initial policy (how it selects actions) and value functions (if applicable).
- Initialize parameters for learning, such as learning rate and exploration strategy (e.g., epsilon-greedy).





# Steps in Reinforcement Learning

Reinforcement Learning (RL) typically follows these key steps:

## **3. Agent-Environment Interaction:**

- The agent observes the current state of the environment.
- Based on its policy, the agent selects an action.
- The environment responds with a reward and the next state.

## **4. Update the Policy and/or Value Function:**

- Use the reward and the observed state transition to update the agent's knowledge.
- Methods include:
  - Value-based updates (e.g., Q-Learning).
  - Policy-based updates (e.g., Policy Gradient).
  - Combined updates (e.g., Actor-Critic).



# Steps in Reinforcement Learning


Reinforcement Learning (RL) typically follows these key steps:

## **5. Balance Exploration and Exploitation:**

- Adjust the exploration strategy to ensure a good balance:
- Explore new actions to find potentially better rewards.
- Exploit known actions to maximize cumulative rewards.

## **6. Repeat Until Convergence:**

- Repeat the agent-environment interaction and updates until the policy converges or performance stabilizes.
- Monitor performance metrics such as total reward or accuracy.



# Steps in Reinforcement Learning


Reinforcement Learning (RL) typically follows these key steps:

## **7. Test the Learned Policy:**

- Evaluate the agent in the environment to ensure it has learned the desired behavior.
- Adjust parameters or re-train if needed.

## **8. Optimize and Deploy:**

- Fine-tune the model for efficiency.
- Deploy the agent in the real-world environment or integrate it into the application.



# Applications of Reinforcement Learning


Reinforcement Learning (RL) has diverse applications across industries and research areas due to its ability to optimize sequential decision-making problems. Here are notable applications:

## 1. Gaming:

- ***Video Games:*** Training AI to play games like Atari, Dota 2, and StarCraft.
- ***Board Games:*** Solving complex games like Chess and Go (e.g., AlphaGo).
- ***Game Testing:*** Automating and optimizing game testing processes.

## 2. Robotics:

- ***Robot Control:*** Teaching robots to walk, run, or perform tasks.
- ***Industrial Automation:*** Optimizing robotic arms for manufacturing and assembly lines.
- ***Drone Navigation:*** Training drones for autonomous flight and obstacle avoidance.



# Applications of Reinforcement Learning


Here are notable applications:

## 3. Healthcare:

- ***Treatment Plans:*** Personalized medical treatment (e.g., optimizing chemotherapy schedules).
- ***Drug Discovery:*** Identifying optimal chemical reactions and compounds.
- ***Healthcare Operations:*** Streamlining patient flow in hospitals.

## 4. Autonomous Vehicles:

- ***Self-Driving Cars:*** Training cars to navigate traffic, obey rules, and make real-time decisions.
- ***Fleet Optimization:*** Routing and scheduling of autonomous delivery vehicles.





# Applications of Reinforcement Learning


Here are notable applications:

## 5. Finance:

- ***Portfolio Optimization:*** Maximizing returns by balancing risk and rewards in investments.
- ***Algorithmic Trading:*** Developing RL-based trading bots for financial markets.
- ***Fraud Detection:*** Dynamic adaptation to detect and prevent fraudulent activities.

## 6. Energy Systems:

- ***Grid Management:*** Optimizing energy distribution and consumption in smart grids.
- ***Renewable Energy:*** Managing resources like solar or wind energy for efficiency.
- ***HVAC Systems:*** Controlling heating and cooling systems in buildings to save energy.



# Applications of Reinforcement Learning


Here are notable applications:

## 7. Natural Language Processing (NLP):

- ***Conversational AI:*** Improving chatbots and virtual assistants for better dialogue handling.
- ***Text Summarization:*** Using RL for coherent and concise summaries.
- ***Language Translation:*** Enhancing the accuracy of machine translation models.

## 8. E-commerce and Marketing:

- ***Dynamic Pricing:*** Optimizing product prices based on demand and competition.
- ***Recommendation Systems:*** Personalizing product or content suggestions for users.
- ***Ad Placement:*** Optimizing online advertisements to maximize revenue.



# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 1. Sample Efficiency:

- **Problem:** RL algorithms often require a large number of interactions with the environment to learn effectively.
- **Impact:** Makes RL impractical in real-world scenarios where interactions are costly or time-consuming (e.g., robotics, healthcare).

## 2. Exploration vs. Exploitation Trade-off:

- **Problem:** Balancing the need to explore new actions with exploiting known actions for rewards is non-trivial.
- **Impact:** Poor exploration can lead to suboptimal policies; excessive exploration wastes resources.



# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 3. Sparse and Delayed Rewards:

- **Problem:** Many environments provide rewards infrequently or after long sequences of actions.
- **Impact:** Makes it difficult for the agent to learn which actions are beneficial.

## 4. High-Dimensional State and Action Spaces:

- **Problem:** Real-world problems often involve complex environments with numerous states and actions.
- **Impact:** Increases computational and memory requirements, making learning inefficient or infeasible.



# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 5. Non-Stationary Environments:

- **Problem:** The environment may change over time (e.g., user behavior in recommendation systems).
- **Impact:** Requires the agent to adapt continuously, complicating policy learning.

## 6. Partial Observability:

- **Problem:** The agent may not have full access to the environment's state (e.g., noisy or missing data).
- **Impact:** Requires advanced techniques like recurrent models or belief-state tracking.





# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 7. Reward Engineering:

- **Problem:** Designing appropriate reward functions is complex and critical to successful learning.
- **Impact:** Poorly designed rewards can lead to unintended behaviors or failed training.

## 8. Scalability:

- **Problem:** RL algorithms struggle to scale efficiently in large or multi-agent environments.
- **Impact:** Limits applicability in domains like traffic control or multi-robot systems.



# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 9. Safety and Ethical Concerns:

- **Problem:** Ensuring safe exploration and decision-making in critical systems (e.g., healthcare, autonomous vehicles).
- **Impact:** Mistakes during learning can have catastrophic consequences.

## 10. Generalization:

- **Problem:** RL agents often overfit to the training environment and fail to generalize to new, unseen scenarios.
- **Impact:** Reduces the robustness and reliability of RL systems in real-world deployments.



# Challenges in Reinforcement Learning


Here are some key challenges in Reinforcement Learning (RL) summarized:

## 11. Computation Costs:

- ***Problem:*** Training deep RL models demands substantial computational resources.
- ***Impact:*** Limits accessibility to individuals or organizations without advanced hardware.

## 12. Multi-Agent Coordination:

- ***Problem:*** Multi-agent RL requires agents to collaborate or compete effectively.
- ***Impact:*** Increases complexity due to dynamic and potentially adversarial interactions.



## • *Example:* FAQs AI Automator

```

1 import os
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.pipeline import make_pipeline
5 from sklearn.feature_extraction.text import CountVectorizer
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import accuracy_score
8
9 # Function to load FAQs from the CSV file
10 def load_faqs(csv_path):
11     try:
12         data = pd.read_csv(csv_path)
13         if 'Question' not in data.columns or 'Response' not in data.columns:
14             raise ValueError("CSV file must contain 'Question' and 'Response' columns.")
15         return data
16     except Exception as e:
17         print(f"Error loading FAQs: {e}")
18         return pd.DataFrame(columns=['Question', 'Response'])
19
20 # Function to display and interact with FAQs
21 def show_faqs(data):
22     if data.empty:
23         print("No FAQs available.")
24         return
25     print("\nFrequently Asked Questions:")
26     for i, row in data.iterrows():
27         print(f"{i+1}. {row['Question']}")
28     print("\nEnter the number of the question you'd like to know more about or type 'back' to return.")
29     while True:
30         user_input = input("Your choice: ").strip()
31         if user_input.lower() == 'back':
32             break
33         try:
34             choice = int(user_input) - 1
35             if 0 <= choice < len(data):
36                 print(f"\n{data.iloc[choice]['Question']}\nBot: {data.iloc[choice]['Response']}\n")
37             else:
38                 print("Invalid choice. Please try again.")
39         except ValueError:
40             print("Invalid input. Please enter a number.")
41

```

- **data.iloc[choice]:** Accesses the specified row in the DataFrame by its integer index (choice). Index-based selection method (**iloc**)
- **['Question'] and ['Response']:** Retrieve the question and response text from the selected row.

- *Example:* FAQs AI Automator

```
42 # Function to evaluate the model
43 def evaluate_model(model, X_test, y_test):
44     try:
45         y_pred = model.predict(X_test)
46         accuracy = accuracy_score(y_test, y_pred)
47         print(f"Model accuracy: {accuracy * 100:.2f}%")
48         return accuracy
49     except Exception as e:
50         print(f"Error during evaluation: {e}")
51         return 0.0
52
53 # Function for interacting with the chatbot
54 def interact_with_chatbot(model, faqs_data):
55     print("\nChatbot is ready! Type 'exit' to quit or 'faq' to check FAQs.")
56     while True:
57         user_input = input("You: ")
58         if user_input.strip().lower() == 'exit':
59             print("Goodbye!")
60             break
61         elif user_input.strip().lower() == 'faq':
62             show_faqs(faqs_data)
63         else:
64             try:
65                 response = model.predict([user_input])[0]
66                 print(f"Bot: {response}")
67             except Exception as e:
68                 print(f"Error generating response: {e}")
69
```



## • *Example:* FAQs AI Automator

```
70 # Main function
71 if __name__ == "__main__":
72     # Path to the CSV file
73     csv_path = "cleaned_data.csv" # Replace with your CSV file path
74
75     # Load FAQs and training data
76     if not os.path.exists(csv_path):
77         print(f"Error: CSV file not found at '{csv_path}'.")
78         exit()
79
80     data = load_faqs(csv_path)
81     if data.empty:
82         print("No valid data found in the CSV file.")
83         exit()
84
85     X = data['Question']
86     y = data['Response']
87
88     # Train-test split
89     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
90
91     # Train the model dynamically
92     print("Training the model...")
93     model = make_pipeline(
94         CountVectorizer(stop_words='english'),
95         LogisticRegression(max_iter=1000, solver='liblinear')
96     )
97     model.fit(X_train, y_train)
98
99     # Evaluate the model
100     accuracy = evaluate_model(model, X_test, y_test)
101     if accuracy < 0.8:
102         print("Warning: Model accuracy is below 80%. Consider improving the dataset.")
103
104     # Interact with the chatbot
105     interact_with_chatbot(model, data)
```

## 4. Homework

- What is reinforcement learning, and how does it differ from supervised and unsupervised learning?
- What are the key characteristics that define reinforcement learning systems?
- What are the two primary types of reinforcement learning, and how do they differ in their approaches?
- What are the main steps involved in a reinforcement learning process?
- What are some of the major challenges faced in reinforcement learning, and how can they be mitigated?





សាកលវិទ្យាល័យនំរតុន  
NORTON UNIVERSITY



# Thank you